

Decision trees

Predictive Analytics



Acknowledgments

These slides draw upon and adapt selected materials by:

- **Fragkiskos D. Malliaros** (CentraleSupélec, Université Paris-Saclay, France)
- **David Sontag** (MIT CSAIL, USA)

Another Classification Idea (2/2)


- Gives axes aligned decision boundaries



Example of a Decision Tree

Tid	Refund	Marital Status	Taxable Income	Cheat	class
1	Yes	Single	125K	No	
2	No	Married	100K	No	
3	No	Single	70K	No	
4	Yes	Married	120K	No	
5	No	Divorced	95K	Yes	
6	No	Married	60K	No	
7	Yes	Divorced	220K	No	
8	No	Single	85K	Yes	
9	No	Married	75K	No	
10	No	Single	90K	Yes	


Training Data



Model: Decision Tree

Another Example of Decision Tree


Tid	Refund	Marital Status	Taxable Income	Cheat	categorical categorical continuous class
1	Yes	Single	125K	No	
2	No	Married	100K	No	
3	No	Single	70K	No	
4	Yes	Married	120K	No	
5	No	Divorced	95K	Yes	
6	No	Married	60K	No	
7	Yes	Divorced	220K	No	
8	No	Single	85K	Yes	
9	No	Married	75K	No	
10	No	Single	90K	Yes	



There could be more than one tree that fits the same data!

Decision Trees – Nodes and Branching

- Internal nodes correspond to **test attributes**
- Branching is determined by attribute value
- Leaf nodes are outputs (class assignments)
 - YES or NO in this example




Decision Tree Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set


Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



Apply Model to Test Data

Start from the root of tree




Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Apply Model to Test Data

Test Data


Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data


Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data


Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data


Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?




Decision Tree Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



Decision Tree Induction – The Idea (1/2)

- Basic algorithm
 - Tree is constructed in a top-down recursive manner
 - Initially, all the training examples are at the root
 - Attributes are categorical (if continuous-valued, they are discretized in advance)
 - Examples are partitioned recursively based on the selected attributes
 - Split attributes are selected on the basis of a heuristic or statistical measure (e.g., gini index, information gain)
- Most commercial DTs use variations of this algorithm

Decision Tree Induction – The Idea (2/2)

- Simple, greedy, recursive approach, builds up tree node-by-node
 1. Pick an attribute to **split** at a non-terminal node
 2. Split examples into groups based on attribute value
 3. For each group:
 - If no examples – return class majority from parent node
 - Else, if all examples are in the same class – return class
 - Else, loop to step 1


Decision Tree Induction Algorithms

- Many algorithms
 - Hunt's algorithm (one of the earliest)
 - CART
 - ID3, C4.5
 - SLIQ, SPRINT


General Structure of Hunt's Algorithm

- Let D_t be the set of training records that reach a node t
- General procedure:
 - If D_t contains records that belong to the same class y_t , then t is a leaf node labeled as y_t
 - If D_t is an empty set, then t is a leaf node labeled by the default class, y_d
 - If D_t contains records that belong to more than one classes, use an attribute test to **split** the data into smaller subsets
 - Recursively apply the procedure to each subset

Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes




Hunt's Algorithm




Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Hunt's Algorithm



Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Hunt's Algorithm



Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Tree Induction


- Greedy strategy
 - Split the records based on an attribute test that optimizes a certain criterion
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition?
 - How to determine the best split?
 - Determine when to stop splitting

How to Specify Test Condition?

- Depends on attribute types
 - Nominal (categorical)
 - E.g., female, male
 - Ordinal (categorical, but there is an ordering)
 - E.g., economic status: low, medium, high
 - Continuous
- Depends on the number of ways to split
 - 2-way split
 - Multi-way split


Splitting Based on Nominal Attributes

- Multi-way split: Use as many partitions as distinct values




No ordering of the attributes

- Binary split: Divides values into two subsets
(Need to find optimal partitioning)




OR




Splitting Based on Ordinal Attributes


- **Multi-way split:** Use as many partitions as distinct values



- **Binary split:** Divides values into two subsets
(Need to find optimal partitioning)




OR




Splitting Based on Continuous Attributes

- Different ways of handling
 - Discretization to form an ordinal categorical attribute
 - Static – discretize once at the beginning
 - Dynamic – ranges can be found by equal interval bucketing, equal frequency bucketing (percentiles), or clustering.
 - Binary decision: $(A < v)$ or $(A \geq v)$
 - Considers all possible splits and finds the best cut
 - Can be more computational intensive

Splitting Based on Continuous Attributes



(i) Binary split



(ii) Multi-way split


Tree Induction

- Greedy strategy
 - Split the records based on an attribute test that optimizes certain criterion
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition?
 - How to determine the best split?
 - Determine when to stop splitting

How to Determine the Best Split (1/2)

Before splitting: 10 records of class 0
10 records of class 1

Suppose that we want to predict if a student will get tax return or no?




Which test condition is the best?

How to Determine the Best Split (1/2)

Before splitting: 10 records of class 0
10 records of class 1

Suppose that we want to predict if a student will get tax return or no?



Which test condition is the best?

How to Determine the Best Split (2/2)


- Greedy approach:
 - Nodes with **homogeneous (pure)** class distribution are preferred
- Need a measure of node impurity:

C0: 5
C1: 5


Non-homogeneous,
High degree of impurity

C0: 9
C1: 1


Homogeneous,
Low degree of impurity



Very impure



Medium impure



Pure

Measures of Node Impurity



- Gini Index
- Entropy and Information Gain
- Misclassification error

How to Find the Best Split – Gini Index

Before Splitting:

C0	N00
C1	N01

Two possible splits
on attributes A or B



C0	N10
C1	N11

C0	N20
C1	N21

C0	N30
C1	N31

C0	N40
C1	N41



How to Find the Best Split – Gini Index

Before Splitting:

C0	N00
C1	N01

M0

Two possible splits
on attributes A or B



C0	N10
C1	N11

M1

C0	N20
C1	N21

M2

C0	N30
C1	N31

M3

C0	N40
C1	N41

M4



How to Find the Best Split – Gini Index

Before Splitting:

C0	N00
C1	N01

M0

Two possible splits
on attributes A or B



C0	N10
C1	N11

M1

C0	N20
C1	N21

M2

C0	N30
C1	N31

M3


C0	N40
C1	N41

M4

M12

M34

How to Find the Best Split – Gini Index



$$\text{Gain} = M_0 - M_{12} \text{ vs. } M_0 - M_{34}$$

Reduction in impurity

Measure of Impurity: Gini Index

Gini index for a node t

$$\text{Gini}(t) = 1 - \sum_{j=1}^J (p(j|t))^2$$

- J number of classes
- $p(j|t)$ is the relative frequency of class j at node t

C1	0
C2	6

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$\text{Gini} = 1 - P(C1)^2 - P(C2)^2 = 1 - 0 - 1 = 0$$

C1	1
C2	5

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$\text{Gini} = 1 - (1/6)^2 - (5/6)^2 = 0.278$$

C1	2
C2	4

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$\text{Gini} = 1 - (2/6)^2 - (4/6)^2 = 0.444$$

Select the split with the smallest Gini index

Splitting based on Gini Index

- When a node t is split on attribute A into k partitions (children), the quality of split is computed as

$$\text{Gini}_A = \sum_{i=1}^k \frac{n_i}{n} \text{Gini}(i) \quad n_i = \text{number of records at child } i \\ n = \text{number of records at node } t$$


- Reduction in impurity (gain) after splitting node t (on attribute A), is computed as

$$\Delta \text{Gini}_A = \text{Gini}(t) - \text{Gini}_A$$

- The attribute X with the smallest Gini_X or the largest reduction in impurity is chosen to split the node

Example

- Splits into two partitions
 - Goal: reduce impurity



$\text{Gini}(N_1)$

$$= 1 - (5/7)^2 - (2/7)^2 \\ = 0.408$$

$\text{Gini}(N_2)$

$$= 1 - (1/5)^2 - (4/5)^2 \\ = 0.32$$

	Parent
C1	6
C2	6
Gini = 0.500	

	N1	N2
C1	5	1
C2	2	4
Gini=0.371		

Split on attribute B

$$\text{Gini}_B = \frac{7}{12} * 0.408 + \frac{5}{12} * 0.32 \\ = 0.371$$

Tree Induction

- Greedy strategy
 - Split the records based on an attribute test that optimizes certain criterion
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition?
 - How to determine the best split?
 - Determine when to stop splitting

Stopping Criteria for Tree Induction

- Stop expanding a node when all the records belong to the same class
- Stop expanding a node when all the records have similar attribute values

Decision Tree Based Classification

- Advantages
 - Inexpensive to construct (training phase)
 - Extremely fast at testing phase (classifying unseen data)
 - **Easy to interpret** for small-sized trees
 - Accuracy is comparable to other classification techniques for many simple data sets

Overfitting and Tree Pruning

- **Overfitting:** An induced tree may overfit the training data
 - Too many branches, some may reflect anomalies due to noise or outliers
- Two approaches to avoid overfitting
 - **Pre-pruning:** Halt tree construction early; do not split a node if this would result in the goodness measure falling below a threshold
 - Difficult to choose an appropriate threshold
 - **Post-pruning:** Remove branches from a fully grown tree
 - Get a sequence of progressively pruned trees
 - Use a dataset (different from the training data) to decide which is the best pruned tree

scikit-learn




[http://scikit-learn.org/stable/modules/generated/
sklearn.tree.DecisionTreeClassifier.html#sklearn.tree.DecisionTreeClassifier](http://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html#sklearn.tree.DecisionTreeClassifier)

Ensemble learning

Ensemble Methods

- Typical application: classification
- Ensemble of classifiers: **set of classifiers** whose individual decisions are combined in some way to classify new examples
- Simplest approach:
 1. Generate multiple classifiers (e.g., decision trees, logistic regression)
 2. Each classifier votes (decides) on a test instance
 3. Take majority as classification
- Classifiers are different due to different sampling of training data, or randomized parameters within the classification algorithm
- **Goal:** take a simple algorithm and transform it into a ‘super classifier’


Ensemble Learning – General Idea



Ensemble Methods: Summary

- Differ in **training strategy** and **combination method**
- **Bagging (bootstrap aggregation)**
 - Random sampling with replacement
 - Train separate models on overlapping training sets, average their predictions
 - E.g., **random forest classifier**
- **Boosting**
 - Sequential training, iteratively re-weighting the training examples – the current classifier focuses on hard examples
 - E.g., **AdaBoost**

Bagging vs. Boosting



Source: <https://quantdare.com/what-is-the-difference-between-bagging-and-boosting/>

Bagging: Bootstrap Estimation

- Repeatedly draw n samples from D
- For each set of samples, estimate a statistic
- The bootstrap estimate is the mean of the individual estimates
- Used to estimate a statistic (parameter) and its variance
- Bagging: bootstrap aggregation (Breiman 1994)

Bagging

- Simple idea
 - Generate **M bootstrap samples** from your original training set
 - Train on each one to **get y_m** , and average them

$$y_{bag}^M(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M y_m(\mathbf{x})$$

- Each bootstrap sample is drawn with **replacement**
 - Each one contains some duplicates of certain training points and leaves out other training points completely
- **For regression:** average the predictions
- **For classification:** average class probabilities or take majority vote

Boosting

- Also works by manipulating the training set, but **the individual classifiers are trained sequentially**
- Each classifier is trained based on knowledge of the performance of previously trained classifiers
 - Focus on difficult examples
- Final classifier: weighted sum of the component classifiers

AdaBoost: Making Weak Learners Stronger

- Suppose you have a weak learning module (a **base classifier**) that can always get $0.5 + \epsilon$ correct when given a binary classification task
 - A little bit better than a random classifier
 - Weak learner
- Can you apply this learning module many times to get a strong learner that can get close to zero error rate on the training data?
 - ML theorists showed how to do this and it actually led to an effective new learning procedure (Freund & Shapire, 1996)
 - The **AdaBoost** algorithm

AdaBoost – The Idea

- Train T weak learners (models) sequentially
- First train the base classifier on all the training data with **equal importance weights** on each case
- Then, re-weight the training data to emphasize the **hard cases** and train a second model
 - Instances that were **misclassified** in the previous step
 - **Q:** How do we re-weight the data?
- Keep training new models on the re-weighted data
- Finally, use a weighted committee of all the models for the test data
 - How do we weight the models in the committee?

How to Train Each Classifier

- Input: $(D_n = \{(\mathbf{x}_i, t_i)\}_{i=1}^n)$
- Output: $y(\mathbf{x}_i) \in \{-1, 1\}$
- Weight of instance (e.g., data point) \mathbf{x}_i for classifier t : $w_i^{(t)}$
- Cost function for classifier t :

$$J_t = \sum_{i=1}^n w_i^{(t)} \underbrace{I(y_t(\mathbf{x}_i) \neq t_i)}_{\text{1 if error, 0 otherwise}} = \sum \text{weighted errors}$$

Weight of Instances for Classifier t


- Weighted error rate of a classifier t

$$\gamma_t = \frac{J_t}{\sum_{i=1}^n w_i^{(t)}}$$

- The quality (coefficient) of classifier t is

It is zero if the classifier has weighted error rate of 0.5 and infinity if the classifier is perfect

$$\alpha_t \leftarrow \ln \left\{ \frac{1 - \gamma_t}{\gamma_t} \right\}$$




- Update the weights for the next classifier ($t+1$)

$$w_i^{(t+1)} \leftarrow w_i^{(t)} \exp \left\{ \alpha_t I(y_t(\mathbf{x}_i) \neq t_i) \right\}$$

The weights 'inform' the training of the weak learner
(decision trees can be grown that favor splitting sets of samples with high weights)

1 if error, 0 otherwise

How to Make Predictions



- Weight the binary prediction of each classifier by the quality of that classifier

$$Y_T(\mathbf{x}) = \text{sign}\left(\sum_{t=1}^T \alpha_t y_t(\mathbf{x})\right)$$

scikit-learn




[http://scikit-learn.org/stable/modules/generated/
sklearn.ensemble.AdaBoostClassifier.html](http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html)

<http://scikit-learn.org/stable/modules/ensemble.html>

Hypotheses: decision trees $f : X \rightarrow Y$

- Each internal node tests an attribute x_i
- One branch for each possible attribute value $x_i=v$
- Each leaf assigns a class y
- To classify input x : traverse the tree from root to leaf, output the labeled y




Human interpretable!

Hypothesis space


- How many possible hypotheses?
- What functions can be represented?

mpg	cylinders	displacement	horsepower	weight	acceleration	modelyear	maker
good	4	low	low	low	high	75to78	asia
bad	6	medium	medium	medium	medium	70to74	america
bad	4	medium	medium	medium	low	75to78	europe
bad	8	high	high	high	low	70to74	america
bad	6	medium	medium	medium	medium	70to74	america
bad	4	low	medium	low	medium	70to74	asia
bad	4	low	medium	low	low	70to74	asia
bad	8	high	high	high	low	75to78	america
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
bad	8	high	high	high	low	70to74	america
good	8	high	medium	high	high	75to83	america
bad	8	high	high	high	low	75to78	america
good	4	low	low	low	low	75to83	america
bad	6	medium	medium	medium	high	75to78	america
good	4	medium	low	low	low	75to83	america
good	4	low	low	medium	high	75to83	america
bad	8	high	high	high	low	70to74	america
good	4	low	medium	low	medium	75to78	europe
bad	5	medium	medium	medium	medium	75to78	europe




What functions can be represented?

- Decision trees can represent any function of the input attributes!
- For Boolean functions, path to leaf gives truth table row
- Could require exponentially many nodes



(Figure from Stuart Russell)





$$cyl=3 \vee (cyl=4 \wedge (maker=asia \vee maker=europe)) \vee \dots$$

Learning *simplest* decision tree is NP-hard

- Learning the simplest (smallest) decision tree is an NP-complete problem [Hyafil & Rivest '76]
- Resort to a greedy heuristic:
 - Start from empty decision tree
 - Split on **next best attribute (feature)**
 - Recurse

Splitting: choosing a good attribute

Would we prefer to split on X_1 or X_2 ?



X_1	X_2	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F
F	T	F
F	F	F

Idea: use counts at leaves to define probability distributions, so we can measure uncertainty!

Entropy


Entropy $H(Y)$ of a random variable Y

$$H(Y) = - \sum_{i=1}^k P(Y = y_i) \log_2 P(Y = y_i)$$

More uncertainty, more entropy!

Information Theory interpretation:

$H(Y)$ is the expected number of bits needed to encode a randomly drawn value of Y (under most efficient code)



High, Low Entropy

- “High Entropy”
 - Y is from a uniform like distribution
 - Flat histogram
 - Values sampled from it are less predictable
- “Low Entropy”
 - Y is from a varied (peaks and valleys) distribution
 - Histogram has many lows and highs
 - Values sampled from it are more predictable

(Slide from Vibhav Gogate)


Entropy Example

$$H(Y) = - \sum_{i=1}^k P(Y = y_i) \log_2 P(Y = y_i)$$

$$P(Y=t) = 5/6$$

$$P(Y=f) = 1/6$$

$$\begin{aligned} H(Y) &= - 5/6 \log_2 5/6 - 1/6 \log_2 1/6 \\ &= 0.65 \end{aligned}$$



X_1	X_2	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F

Conditional Entropy


Conditional Entropy $H(Y|X)$ of a random variable Y conditioned on a random variable X

$$H(Y | X) = - \sum_{j=1}^v P(X = x_j) \sum_{i=1}^k P(Y = y_i | X = x_j) \log_2 P(Y = y_i | X = x_j)$$

Example:

$$P(X_1=t) = 4/6$$

$$P(X_1=f) = 2/6$$



$$\begin{aligned} H(Y|X_1) &= - \frac{4}{6} (1 \log_2 1 + 0 \log_2 0) \\ &\quad - \frac{2}{6} (1/2 \log_2 1/2 + 1/2 \log_2 1/2) \\ &= 2/6 \end{aligned}$$

X_1	X_2	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F

Information gain

- Decrease in entropy (uncertainty) after splitting

$$IG(X) = H(Y) - H(Y | X)$$

In our running example:

$$\begin{aligned} IG(X_1) &= H(Y) - H(Y|X_1) \\ &= 0.65 - 0.33 \end{aligned}$$


$IG(X_1) > 0 \rightarrow$ we prefer the split!

X_1	X_2	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F

Learning decision trees

- Start from empty decision tree
- Split on **next best attribute (feature)**
 - Use, for example, information gain to select attribute:
- Recurse


When to stop?



First split looks good! But, when do we stop?

Base Case One

mpg values: bad good



Base Case Two

mpg values: bad good

