

# Clustering

Predictive Analytics

# Acknowledgments

These slides draw upon and adapt selected materials by:

- **Fragkiskos D. Malliaros** (CentraleSupélec, Université Paris–Saclay, France)

# Unsupervised learning (clustering)

# Supervised vs. Unsupervised Learning

## Supervised learning

- We have labeled examples
- Given those examples, learn a model that can generalize to unseen examples
- Key tasks:
  - Classification
  - Regression

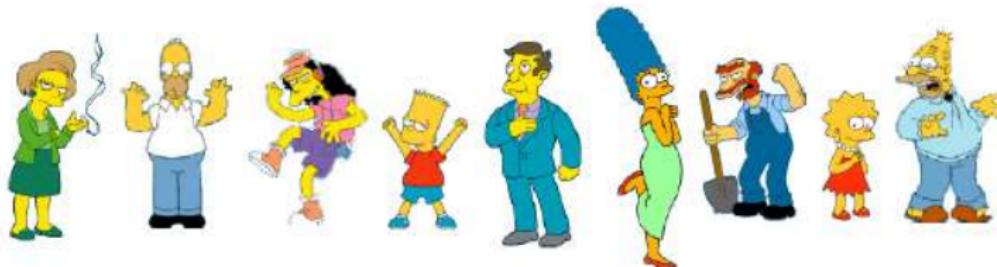
## Unsupervised learning

- The data is unlabeled
- Given the data, learn a model that identifies structure in the data (and generalize to new data)
- Key task:
  - Clustering

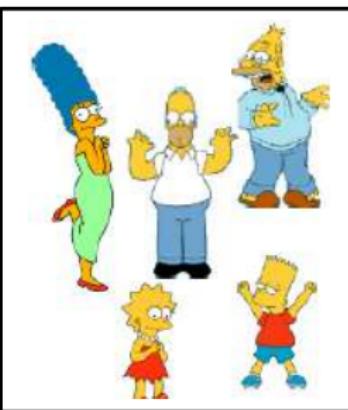
# What is Cluster Analysis?

- Cluster: a collection of data objects
  - **Similar** (or related) to one another within the same group
  - **Dissimilar** (or unrelated) with objects in other groups
- Cluster analysis (or clustering)
  - Finding **similarities** between data according to the characteristics of the data and ...
  - ... **grouping** similar data objects into clusters
- Typical applications
  - As a **stand-alone tool** to get further insights about the data
  - As a **preprocessing step** for other algorithms

# Any Natural Grouping?



clustering is subjective



Simpson's family



School employees



Females



Males

# What is a Good Clustering?

- Good clusters have:
  - High **intra-cluster** similarity: **cohesive** within clusters
  - Low **inter-cluster** similarity: **distinctive** between clusters
- The quality of a clustering method depends on
  - The **similarity measure** used by the method
  - Its ability to discover some or all of the hidden patterns

Recall the distance and similarity measures covered in kNN classification

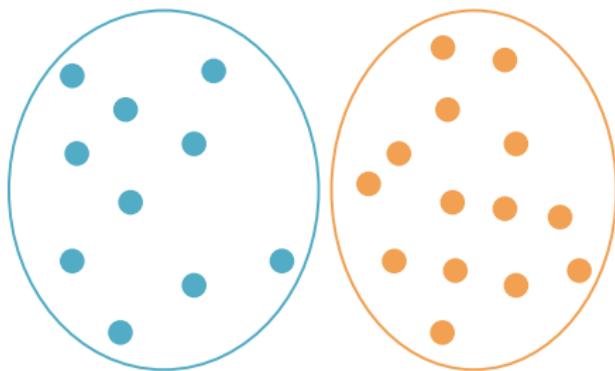
# Goals of Clustering

- Group objects that are similar into **clusters**: classes that are unknown beforehand



# Goals of Clustering

- Group objects that are similar into **clusters**: classes that are unknown beforehand

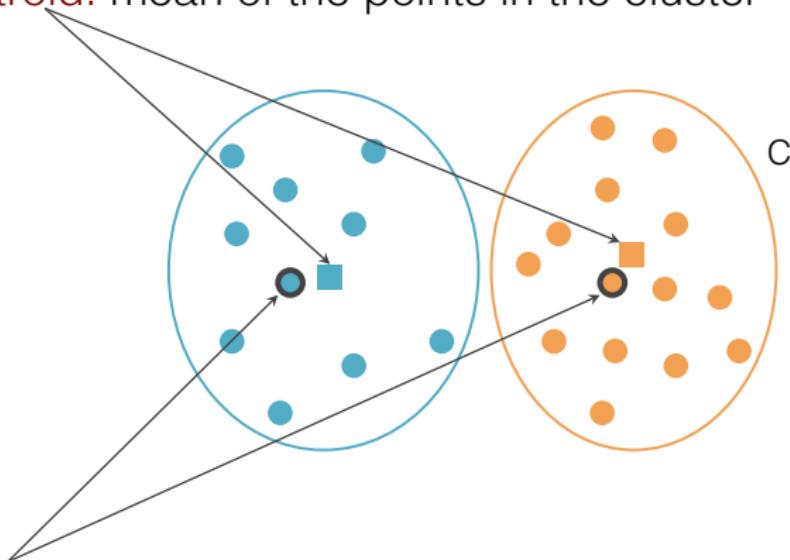


# Applications of Clustering

- Understand general characteristics of the data
- Visualize the data
- Infer some properties of a data point based on how it relates to other data points
- Examples
  - Find subtypes of diseases
  - Visualize protein families
  - Find categories among images
  - Find patterns in financial transactions
  - Detect communities in social networks
  - Find users with similar interests (e.g., Netflix, Amazon)

# Cluster Centroids and Medoids

- **Centroid:** mean of the points in the cluster  $\mu = \frac{1}{|C|} \sum_{x \in C} x$



- **Medoid:** point in the cluster that is closest to the centroid

$$m = \arg \min_{x \in C} d(x, \mu)$$

# Cluster Evaluation

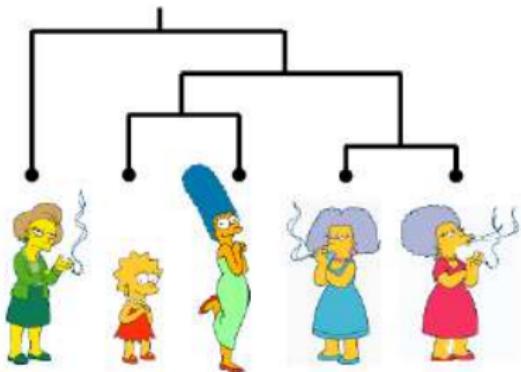
- Clustering is **unsupervised**
- There is no **ground truth**. How do we evaluate the quality of a clustering algorithm?
- Based on the **shape** of the clusters:
  - Points within the same cluster should be nearby/similar and points far from each other should belong to different clusters
- Based on the **stability** of the clusters:
  - We should get the same results if we remove some data points, add noise, etc.
- Based on **domain knowledge**
  - The clusters should “make sense”

# Major Clustering Approaches

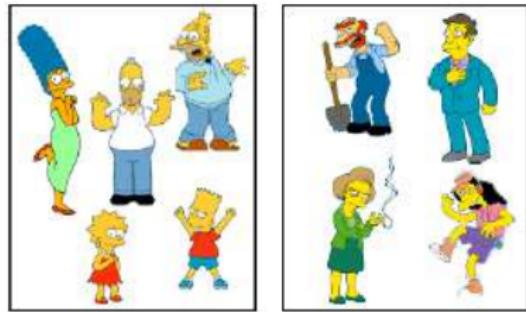
- Partitioning approach
  - Construct various partitions and then evaluate them by some criterion, e.g., minimizing the sum of square errors
  - Typical methods: k-means, k-medoids
- Hierarchical approach
  - Create a hierarchical decomposition of the set of data (or objects) using some criterion
  - Typical methods: Diana, Agnes, BIRCH, CHAMELEON
- Density-based approach
  - Based on connectivity and density functions
  - Typical methods: DBSCAN, OPTICS, DenClue
- Grid-based approach
  - Based on a multiple-level granularity structure
  - Typical methods: STING, WaveCluster, CLIQUE

# Hierarchical vs. Partitional

Hierarchical



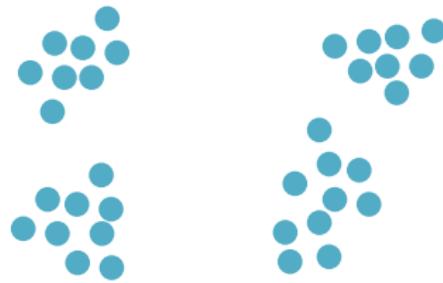
Partitional



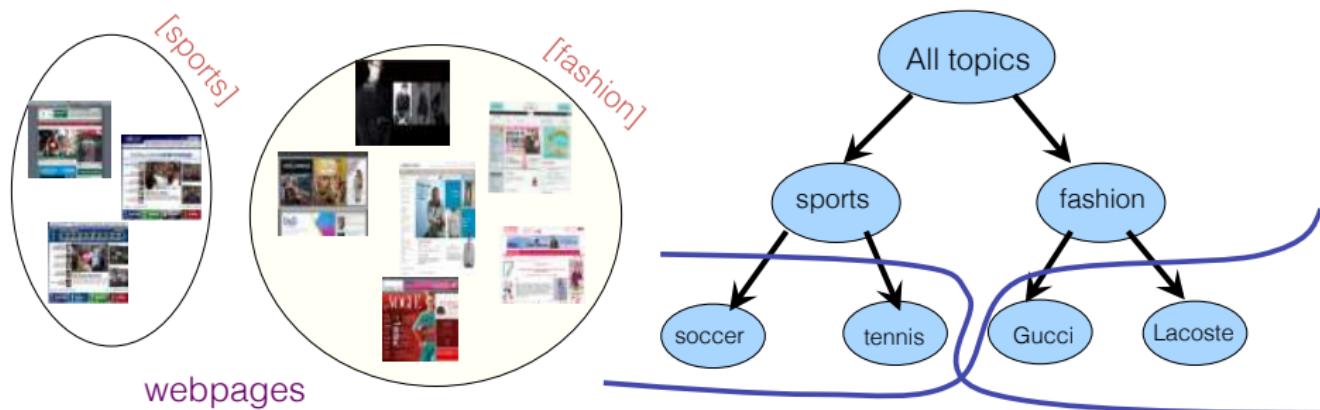
# Hierarchical clustering

# Hierarchical Clustering

- Group data over a variety of possible scales, in a multi-level hierarchy



# Hierarchical Clustering (2/2)



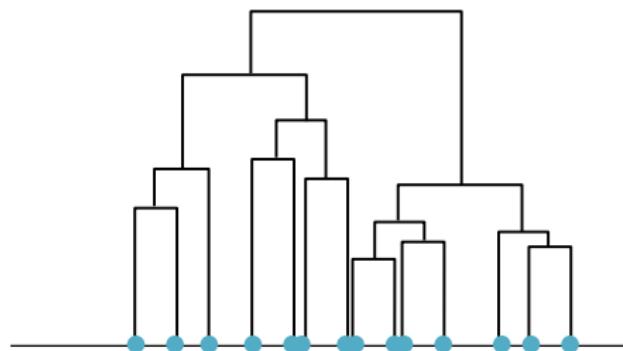
- A hierarchy might be more natural
- Different users might care about different levels of granularity or even prunings

# Construction

- Agglomerative approach (bottom-up)
  - Start with each element in its own cluster
  - Iteratively **join** neighboring clusters
- Divisive approach (top-down)
  - Start with all elements in the same cluster
  - Iteratively **separate** into smaller clusters

# Dendrogram

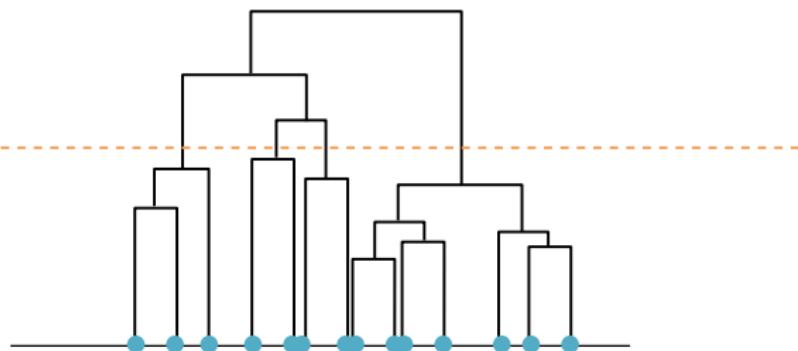
- The results of a hierarchical clustering algorithm are presented in a **dendrogram**



# Dendrogram

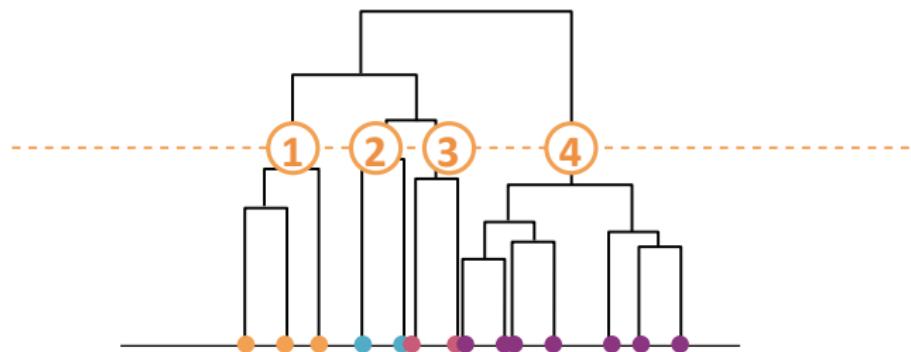
- The results of a hierarchical clustering algorithm are presented in a **dendrogram**

How many clusters  
do I have?



# Dendrogram

- The results of a hierarchical clustering algorithm are presented in a **dendrogram**



# Hierarchical Clustering

- Advantages
  - No need to pre-define the number of clusters
  - Interpretability
- Drawbacks
  - Computational complexity
  - Must decide at which level of the hierarchy to split
  - Lack of robustness (unstable)

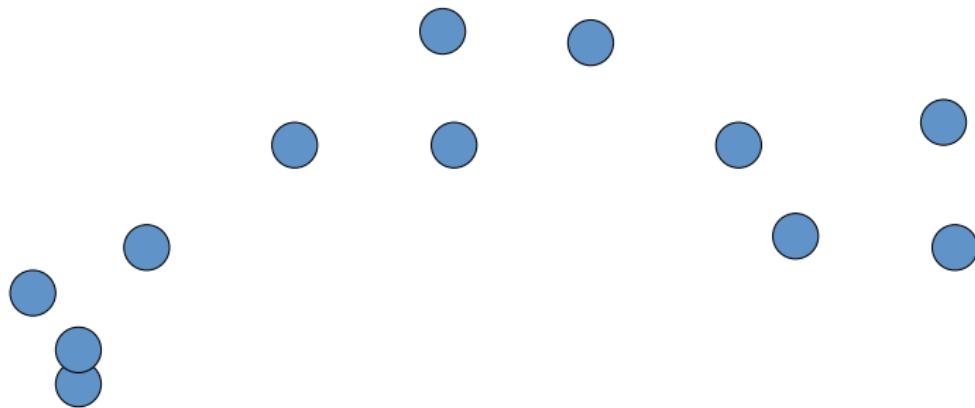
# k-means clustering

# k-means Algorithm – The Idea

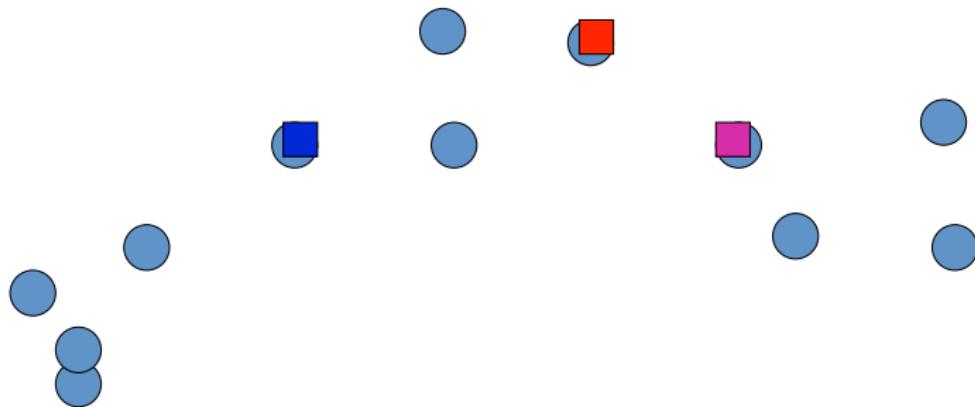
Most well-known and popular clustering algorithm:

1. Start with some initial cluster centers
2. Iterate:
  - Assign each example to closest center
  - Recalculate centers as the mean of the points in a cluster

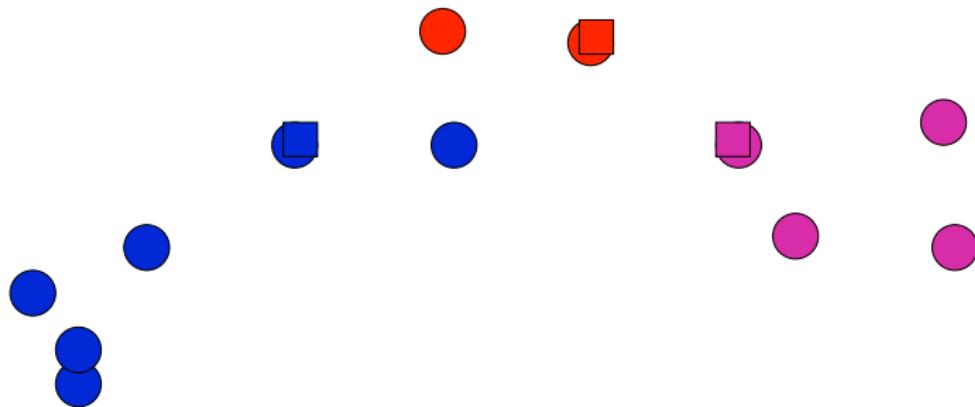
# k-means: An Example



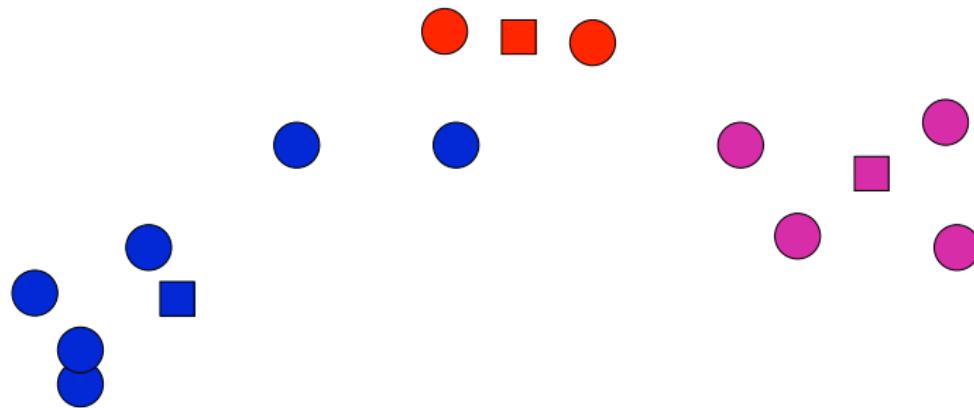
# k-means: Initialize Centers Randomly



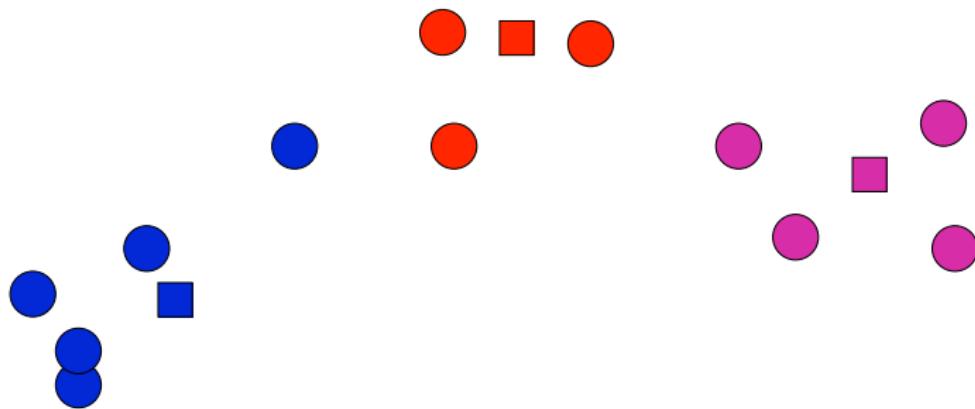
# k-means: Assign Points to Nearest Center



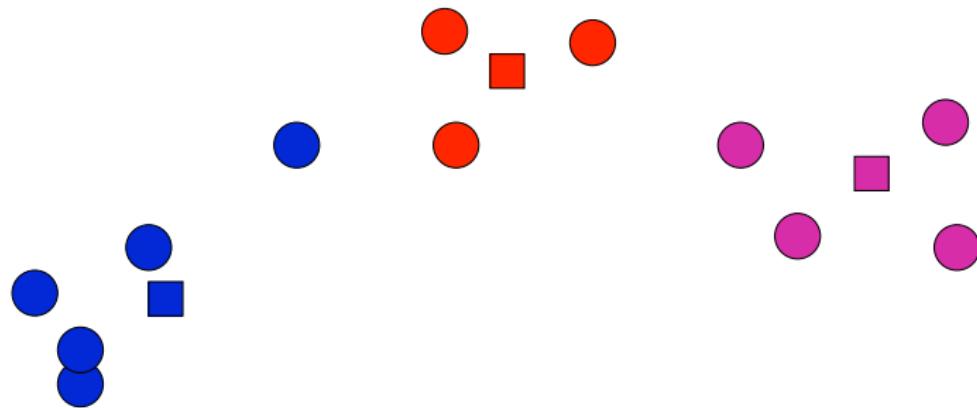
# k-means: Readjust Centers



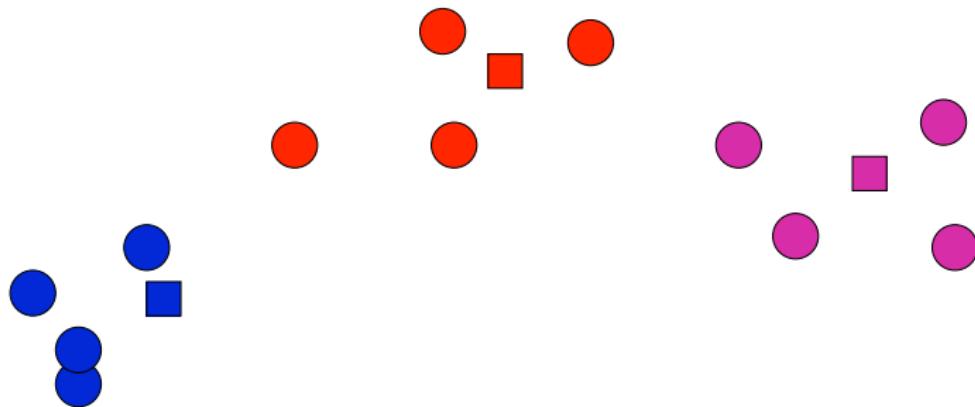
# k-means: Assign Points to Nearest Center



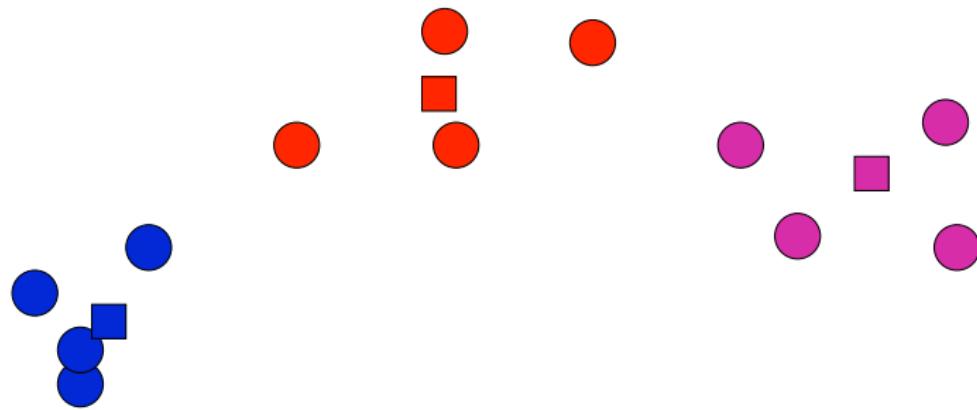
# k-means: Readjust Centers



# k-means: Assign Points to Nearest Center

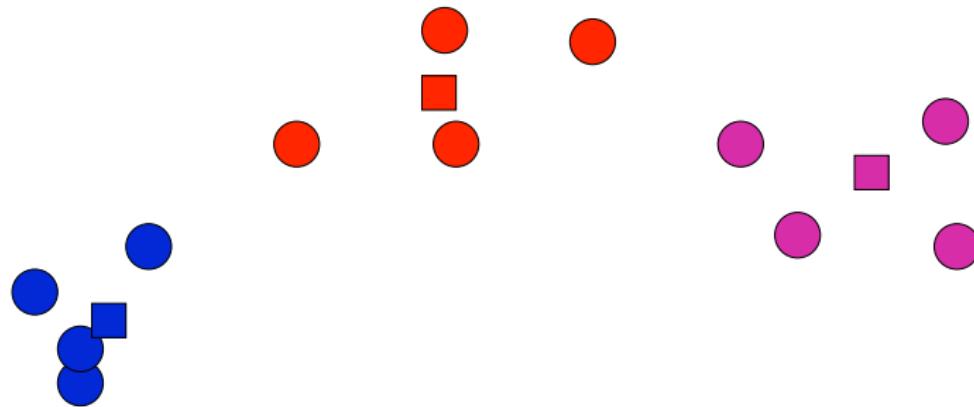


# k-means: Readjust Centers



# k-means: Assign Points to nearest Center

No changes: Done



Test demo at: <http://stanford.edu/class/ee103/visualizations/kmeans/kmeans.html>

# k-means Clustering – Objective Function

- Minimize the intra-cluster variance
  - Within-cluster sum of squares

$$\text{Var}_{\text{in}}(C) = \frac{1}{|C|} \sum_{x \in C} \|x - \mu_C\|^2 \quad \text{For a cluster } C$$

$$V = \sum_{k=1}^K \sum_{x \in C_k} \frac{1}{|C_k|} \|x - \mu_{C_k}\|^2 \quad \text{For all clusters}$$

- What will this partition of the space look like?

# k-means Clustering – Objective Function

- Minimize the intra-cluster variance
  - Within-cluster sum of squares

$$\text{Var}_{\text{in}}(C) = \frac{1}{|C|} \sum_{x \in C} \|x - \mu_C\|^2 \quad \text{For a cluster } C$$

$$V = \sum_{k=1}^K \sum_{x \in C_k} \frac{1}{|C_k|} \|x - \mu_{C_k}\|^2 \quad \text{For all clusters}$$

- For each cluster, the points in that cluster are those that are closest to its centroid than to any other centroid

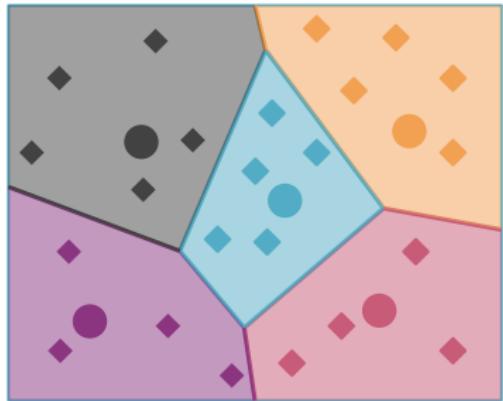
# k-means Clustering – Objective Function

- Minimize the **intra-cluster variance**
  - Within-cluster sum of squares

$$\text{Var}_{\text{in}}(C) = \frac{1}{|C|} \sum_{x \in C} \|x - \mu_C\|^2 \quad \text{For a cluster } C$$

$$V = \sum_{k=1}^K \sum_{x \in C_k} \frac{1}{|C_k|} \|x - \mu_{C_k}\|^2 \quad \text{For all clusters}$$

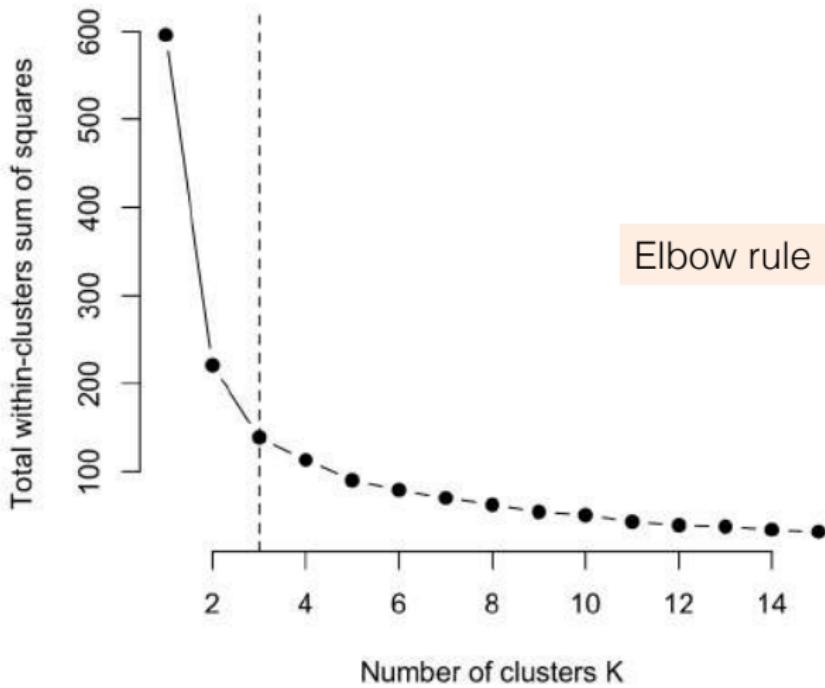
- Voronoi **tesselation**
- Optimal solution: hard problem



# Lloyd's Algorithm for k-means

- k-means cannot be easily optimized
- We adopt a **greedy strategy** (Lloyd's Algorithm)
  - Partition the data into **k** clusters at random
  - Compute the centroid of each cluster
  - Assign each point to the cluster of the centroid centroid
  - Repeat until cluster membership converges

# How to Select k?



# Summary of k-means

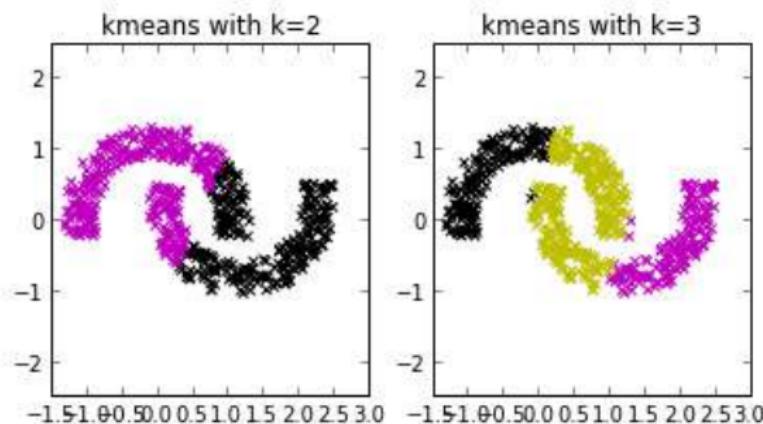
- Advantages

- Computational time is linear:  $\mathcal{O}(npkt)$       t: number of iterations
  - Easily implementable
- ↑  
compute **kn** distances  
in **p** dimensions

- Drawbacks

- Need to select **k** (user-defined parameter)
- Sensitivity to noise and outliers
- Non-deterministic (stochastic)
  - Different solutions with each iteration
- The clusters are forced to have “spherical” (convex) shapes

# Example



Source: <https://pafnuty.wordpress.com/>

# Improving k-means: k-means++

- The quality of the solution depends on the **initialization**
- Rationale behind random initialization
  - Choosing a random assignment may lead the algorithm to a good local minimum
- Another approach: **k-means++** [Arthurs and Vassilvitskii '07]
  1. Select a random point and declare it centroid  $c_1$
  2. For all remaining data points  $x_j$  compute distance  $d(x_j, c_1)$
  3. Select a random point with probability proportional to  $d(x_j, c_1)^2$  and set it as  $c_2$ 
    - This will give a point far enough from  $c_1$
  4. Repeat steps 2 and 3 for all centroids  $k$

# scikit-learn



<http://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>

[http://scikit-learn.org/stable/auto\\_examples/cluster/plot\\_kmeans\\_digits.html#sphx-glr-auto-examples-cluster-plot-kmeans-digits-py](http://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_digits.html#sphx-glr-auto-examples-cluster-plot-kmeans-digits-py)

# scikit-learn

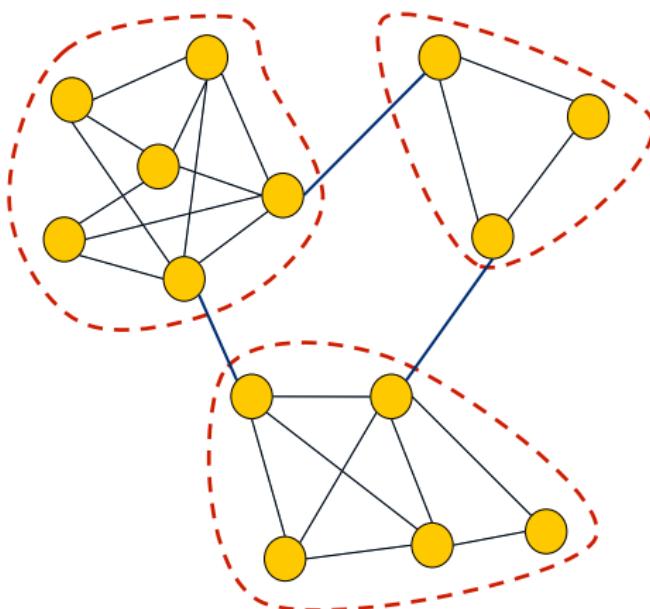


<http://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>

[http://scikit-learn.org/stable/auto\\_examples/cluster/plot\\_kmeans\\_digits.html#sphx-glr-auto-examples-cluster-plot-kmeans-digits-py](http://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_digits.html#sphx-glr-auto-examples-cluster-plot-kmeans-digits-py)

# Spectral clustering

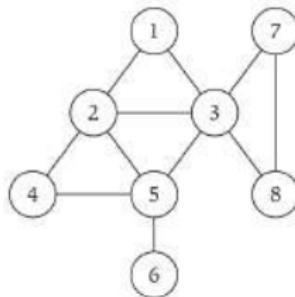
# Clustering Structure in Graphs



How to discover the clustering structure?

# Graph Representation: Adjacency Matrix

- A graph can be represented by the adjacency matrix  $A$ 
  - Matrix of size  $n \times n$ , where  $n = |V|$  is the number of nodes
  - $A_{ij} > 0$ , if  $i$  and  $j$  are connected
  - $A_{ij} = 0$ , if  $i$  and  $j$  are not connected
  - In case of **unweighted** graphs,  $A_{ij} = 1$ , if  $(i, j)$  is an edge of the graph
  - Space proportional to  $n^2$



Undirected graph

Node indexing									
		1	2	3	4	5	6	7	8
Node indexing	1	0	1	1	0	0	0	0	0
	2	1	0	1	1	1	0	0	0
	3	1	1	0	0	1	0	1	1
	4	0	1	0	0	1	0	0	0
	5	0	1	1	1	0	1	0	0
	6	0	0	0	0	1	0	0	0
	7	0	0	1	0	0	0	0	1
	8	0	0	1	0	0	0	1	0

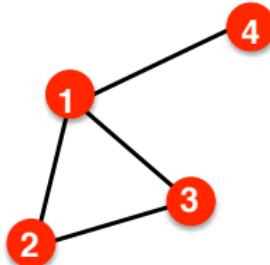
Adjacency matrix

# Graph Representation: Laplacian Matrix

- Let  $\mathbf{G} = (\mathbf{V}, \mathbf{E})$  be a graph. Then, the **Laplacian matrix** is defined as

$$\mathbf{L}_{ij} = \begin{cases} k_i & \text{if } i = j \\ -1 & \text{if } (i, j) \in E \\ 0 & \text{otherwise} \end{cases}$$

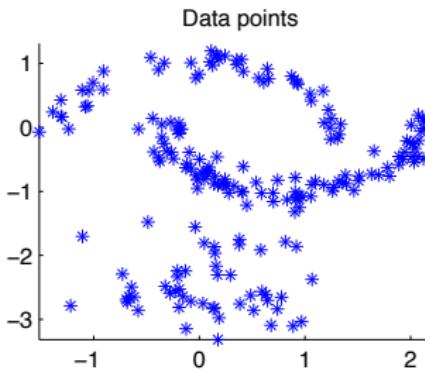
- Diagonal degree matrix  $\mathbf{D}$ , where  $D_{ii} = k_i$  (node degrees)



$$\mathbf{L} = \mathbf{D} - \mathbf{A} = \begin{bmatrix} 3 & -1 & -1 & -1 \\ -1 & 2 & -1 & 0 \\ -1 & -1 & 2 & 0 \\ -1 & 0 & 0 & 1 \end{bmatrix}$$

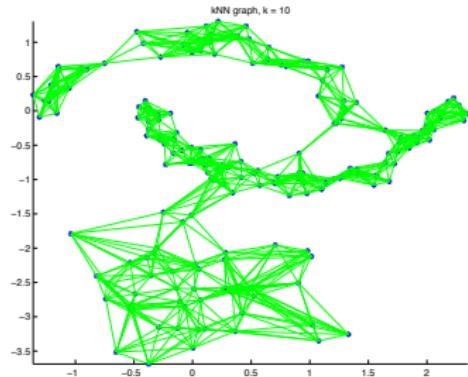
# Clustering Non-Graph Data

- Apply graph clustering algorithms on data with no inherent graph structure (e.g., points in a  $d$ -dimensional Euclidean space)
- How?
  1. Construct a **similarity graph** based on the topological relationships and distances between data points (e.g., kNN graph)
  2. Then, the problem of clustering the set of data points is transformed to a graph clustering problem



Similarity graph  
(e.g., kNN)

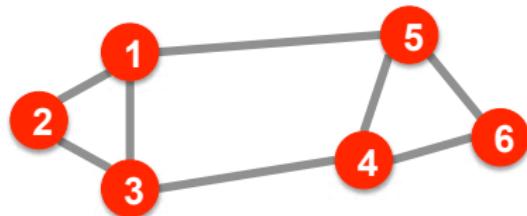
An orange arrow points from the text "Similarity graph (e.g., kNN)" towards the right side of the slide, indicating the transformation process.



[von Luxburg '07], [Shi and Malik '00], [Ng, Jordan, Weiss '02]

# Graph Partitioning (1/2)

- Undirected graph  $G=(V, E)$
- Bi-partitioning task:
  - Divide nodes into two disjoint groups  $A, B$

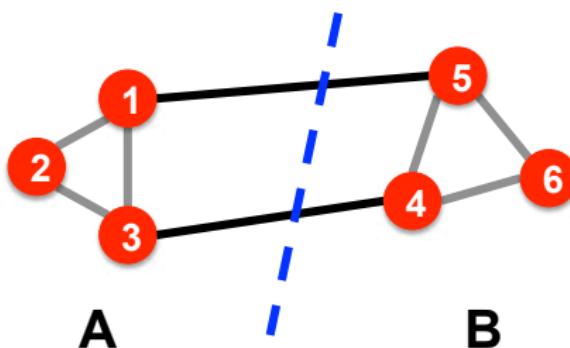


## Questions:

- How can we define a **good** partition of  $G$ ?
- How can we efficiently identify such a partition?

# Graph Partitioning (2/2)

- What makes a **good partition**?
  - Maximize the number of within-group connections
  - Minimize the number of between-group connections

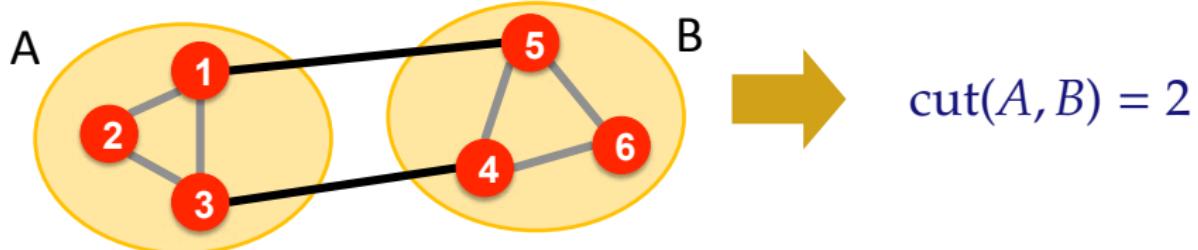


# Graph Cuts

- Express partitioning objectives as a function of the **edge cut** of the partition
- Cut:** Set of edges across two groups:

$$\text{cut}(A, B) = \sum_{i \in A, j \in B} w_{ij}$$

Two partitions, A and B

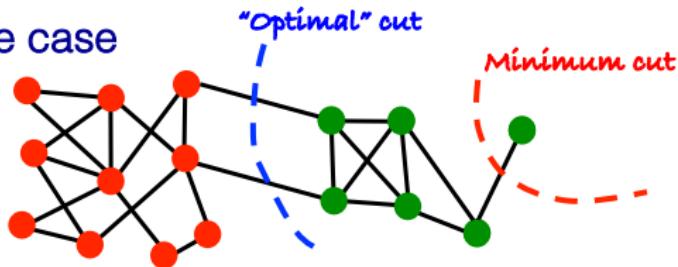


# Graph Cut Criterion for Clustering

- Criterion: Minimum-cut
  - Minimize the weight of connections between groups

$$\arg \min_{A,B} \text{cut}(A,B)$$

- Degenerate case



## Problem

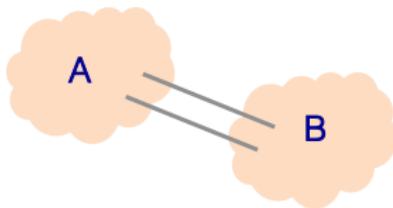
- Not satisfactory partition – often isolated nodes
- Does not consider internal cluster connectivity

# Ratio Cut

Normalize cut by the **size** of the groups

$$\text{ratio-cut}(A, B) = \frac{\text{cut}(A, B)}{|A|} + \frac{\text{cut}(A, B)}{|B|}$$

*Size of A and B*



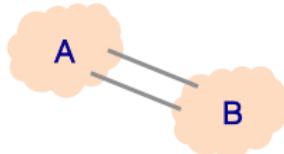
**Internal** group connectivity is not taken into account

# Normalized Cut

- Criterion: **Normalized cut**
  - Connectivity between groups relative to the density of each group

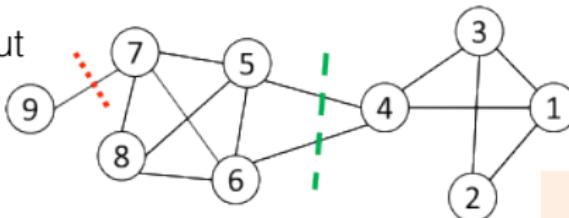
$$\text{normalized-cut}(A, B) = \frac{\text{cut}(A, B)}{\text{vol}(A)} + \frac{\text{cut}(A, B)}{\text{vol}(B)}$$

- Why use this criterion?
  - It produces more **balanced partitions**
- How do we efficiently find a good partition?
  - Computing the optimal cut is **NP-hard**



# Ratio Cut vs. Normalized Cut (1/2)

Red is Min-Cut



$$\text{ratio-cut}(A, B) = \frac{\text{cut}(A, B)}{|A|} + \frac{\text{cut}(A, B)}{|B|}$$

$$\text{normalized-cut}(A, B) = \frac{\text{cut}(A, B)}{\text{vol}(A)} + \frac{\text{cut}(A, B)}{\text{vol}(B)}$$

$$\text{Ratio-Cut(Red)} = 1/1 + 1/8 = 1.125$$

$$\text{Ratio-Cut(Green)} = 2/5 + 2/4 = 0.9$$

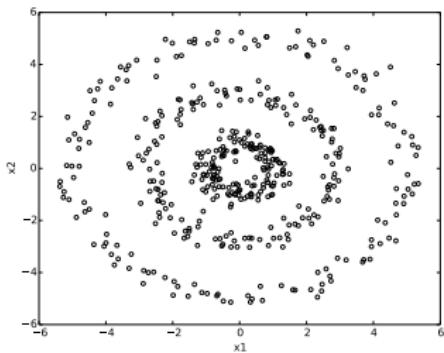
*Lower value is better*

$$\text{Normalized-Cut(Red)} = 1/1 + 1/26 = 1.03$$

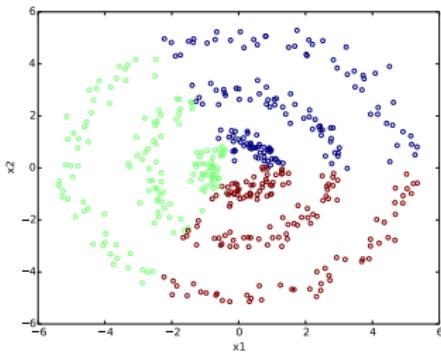
$$\text{Normalized-Cut(Green)} = 2/12 + 2/16 = 0.29$$

Normalized is even better  
for Green due to density

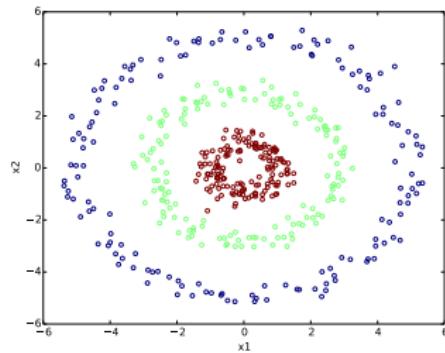
# Spectral Clustering vs. k-Means



- 2-dimensional points
- Find  **$k=3$**  clusters



k-means



spectral clustering

# scikit-learn



<http://scikit-learn.org/stable/modules/clustering.html#spectral-clustering>

[http://scikit-learn.org/stable/auto\\_examples/cluster/  
plot\\_cluster\\_comparison.html](http://scikit-learn.org/stable/auto_examples/cluster/plot_cluster_comparison.html)

# Comparison (scikit-learn)

